

Chinese Ministry of Foreign Affairs Press Conferences Corpus (CMFA PressCon)

Michal Mochtak and Richard Q. Turcsanyi

13 January 2024 (v4.0)

Contents

Overview	1
Citation Information	2
id_orig	2
id_split	2
id	2
link	2
spokesperson	2
day	2
month	2
year	2
date	2
question	2
answer	2
question_lem	2
answer_lem	3
who_asked	3
q_loc	3
q_per	3
q_org	3
q_misc	3
a_loc	3
a_per	3
a_org	3
a_misc	3
q_sentiment	3
a_sentiment	4

Overview

This is a codebook for an original corpus of the Chinese Ministry of Foreign Affairs Press Conferences (CMFA PresCon) mapping two decades of Chinese diplomatic discourse and priorities in China’s foreign policy. The dataset is organized around a question/response structure extracted from the official transcripts of press conferences held between 15 October 2002 and 31 December 2023 and in its current version (v4) counts 30 964 data points.

Citation Information

If you use the dataset, please cite: Mochtak, Michal and Richard Q. Turcsanyi (2021): “Studying Chinese Foreign Policy Narratives: Introducing the Ministry of Foreign Affairs Press Conferences Corpus”. *Journal of Chinese Political Science*, 26 (4): 743-761.

id_orig

Original ID assigned to collected question-response dyad; ordered by date.

id_split

ID assigned to question-response dyad extracted from a multi-topical question. It contains the information on the original question-response ID (X_) and an index (_y) capturing the split.

id

ID assigned to question-answer dyad after the multi-topical split; ordered by date and multi-topical split (in a sequence).

link

Link originally used for accessing the transcript of a press conference (some of them do not work anymore).

spokesperson

Spokesperson holding the press conference.

day

Day of the press conference.

month

Month of the press conference.

year

Year of the press conference.

date

Full date of the press conference.

question

Question(s) asked by a member of the press corps.

answer

Answer(s) provided by the spokesperson.

question_lem

A lemmatized version of “question”. Lemmatization was done using Trankit and XLM-Roberta-base model.

answer_lem

A lemmatized version of “answer”. Lemmatization was done using Trankit and XLM-Roberta-base model.

who_asked

Authorship of the question (news outlet). Period covered: 24 March 2020 - 31 December 2023.

q_loc

Named entities for locations extracted from a transcribed question.

q_per

Named entities for persons extracted from a transcribed question.

q_org

Named entities for organizations extracted from a transcribed question.

q_misc

Named entities for miscellaneous references extracted from a transcribed question.

a_loc

Named entities for locations extracted from a transcribed answer.

a_per

Named entities for persons extracted from a transcribed answer.

a_org

Named entities for organizations extracted from a transcribed answer.

a_misc

Named entities for miscellaneous references extracted from a transcribed answer.

q_sentiment

A sentiment score of a question. The sentiment is an aggregated value of sentiment scores of sentences in a question. The score is obtained through a domain-specific deep learning model trained for the detection of sentiment in political discourse (XLM-R-ParlaSent; <https://huggingface.co/classla/xlm-r-parlasent>). XLM-R-ParlaSent is a multilingual sentiment regression model based on xlm-r-parla, an XLM-R-large model additionally pre-trained on parliamentary proceedings collected by the ParlaMint project (<https://www.clarin.eu/parlamint>). The classifier is a regressor trained on numeric labels 0-5, where zero represents the scale’s negative end, and five represents the positive end. In this context, questions with scores closer to 0 are deemed to have a negative sentiment, and questions with scores closer to 5 have a positive sentiment. As a regressor, the model can also produce scores that are above (>5) or below (<0) the values used during the training. See the accompanying paper for more details: Mochtak, M., Rupnik, P. & Ljubešić, N. (2023). The ParlaSent multilingual training dataset for sentiment identification in parliamentary proceedings. arXiv. <https://arxiv.org/abs/2309.09783>.

a_sentiment

A sentiment score of an answer. The sentiment is an aggregated value of sentiment scores of sentences in an answer. The score is obtained through a domain-specific deep learning model trained for the detection of sentiment in political discourse (XLM-R-ParlaSent; <https://huggingface.co/classla/xlm-r-parlasent>). XLM-R-ParlaSent is a multilingual sentiment regression model based on xlm-r-parla, an XLM-R-large model additionally pre-trained on parliamentary proceedings collected by the ParlaMint project (<https://www.clarin.eu/parlamint>). The classifier is a regressor trained on numeric labels 0-5, where zero represents the scale's negative end, and five represents the positive end. In this context, answers with scores closer to 0 are deemed to have a negative sentiment, and answers with scores closer to 5 have a positive sentiment. As a regressor, the model can also produce scores that are above (>5) or below (<0) the values used during the training. See the accompanying paper for more details: Mochtak, M., Rupnik, P. & Ljubešić, N. (2023). The ParlaSent multilingual training dataset for sentiment identification in parliamentary proceedings. arXiv. <https://arxiv.org/abs/2309.09783>.